

Dataset Link:

<https://www.kaggle.com/datasets/ramjasmaurya/cost-prediction-in-foodmart>

PREDICT COST ON MEDIA CAMPAIGNS IN FOOD MART OF USA .

ON THE BASIS OF 60K CUSTOMERS INCOME ,PRODUCT,PROMOTION AND STORE FEATURES.

ABOUT FOODMART:

Food Mart (CFM) is a chain of convenience stores in the United States. The private company's headquarters are located in Mentor, Ohio, and there are currently approximately 325 stores located in the US. Convenient Food Mart operates on the franchise system.

Food Mart was the nation's third-largest chain of convenience stores as of 1988.

The NASDAQ exchange dropped Convenient Food Mart the same year when the company failed to meet financial reporting requirements.

Carden & Cherry advertised Convenient Food Mart with the Ernest character in the 1980s.

```
In [1]: #Import the necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt,seaborn as sns,plotly.express as px
from sklearn.preprocessing import StandardScaler,MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import accuracy_score,classification_report,confusion_mat
import plotly.graph_objects as go
import warnings
warnings.filterwarnings('ignore')

In [2]: #Load the data using the pandas read function
data=pd.read_csv('/home/vinod/Downloads/media prediction and its cost.csv')
#Print the first five rows of the dataset
data.head().style.background_gradient(cmap='winter')
```

Out[2]:

	food_category	food_department	food_family	store_sales(in millions)	store_cost(in millions)	unit_sales(in millions)	promot
0	Breakfast Foods	Frozen Foods	Food	7.360000	2.723200	4.000000	B
1	Breakfast Foods	Frozen Foods	Food	5.520000	2.594400	3.000000	Cas
2	Breakfast Foods	Frozen Foods	Food	3.680000	1.361600	2.000000	I
3	Breakfast Foods	Frozen Foods	Food	3.680000	1.177600	2.000000	Cas
4	Breakfast Foods	Frozen Foods	Food	4.080000	1.428000	3.000000	Do

In [3]: *#Data information to the dataset*
data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 60428 entries, 0 to 60427
Data columns (total 40 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   food_category                        60428 non-null  object
1   food_department                     60428 non-null  object
2   food_family                         60428 non-null  object
3   store_sales(in millions)            60428 non-null  float64
4   store_cost(in millions)             60428 non-null  float64
5   unit_sales(in millions)             60428 non-null  float64
6   promotion_name                      60428 non-null  object
7   sales_country                      60428 non-null  object
8   marital_status                     60428 non-null  object
9   gender                             60428 non-null  object
10  total_children                     60428 non-null  float64
11  education                         60428 non-null  object
12  member_card                       60428 non-null  object
13  occupation                         60428 non-null  object
14  houseowner                        60428 non-null  object
15  avg_cars_at home(approx)           60428 non-null  float64
16  avg. yearly_income                 60428 non-null  object
17  num_children_at_home               60428 non-null  float64
18  avg_cars_at home(approx).1         60428 non-null  float64
19  brand_name                        60428 non-null  object
20  SRP                               60428 non-null  float64
21  gross_weight                      60428 non-null  float64
22  net_weight                        60428 non-null  float64
23  recyclable_package                60428 non-null  float64
24  low_fat                          60428 non-null  float64
25  units_per_case                    60428 non-null  float64
26  store_type                        60428 non-null  object
27  store_city                       60428 non-null  object
28  store_state                      60428 non-null  object
29  store_sqft                       60428 non-null  float64
30  grocery_sqft                     60428 non-null  float64
31  frozen_sqft                      60428 non-null  float64
32  meat_sqft                       60428 non-null  float64
33  coffee_bar                       60428 non-null  float64
34  video_store                      60428 non-null  float64
35  salad_bar                       60428 non-null  float64
36  prepared_food                    60428 non-null  float64
37  florist                         60428 non-null  float64
38  media_type                       60428 non-null  object
39  cost                            60428 non-null  float64
dtypes: float64(23), object(17)
memory usage: 18.4+ MB
```

Let's check the shape of the dataset

data.shape

```
In [4]: #Statistical analysis of the dataset
data.describe().style.background_gradient(cmap='gist_gray_r')
```

Out[4]:

	store_sales(in millions)	store_cost(in millions)	unit_sales(in millions)	total_children	avg_cars_at home(approx)	num_children_at_ho
count	60428.000000	60428.000000	60428.000000	60428.000000	60428.000000	60428.000
mean	6.541031	2.619460	3.093169	2.533875	2.200271	0.829
std	3.463047	1.453009	0.827677	1.490165	1.109644	1.303
min	0.510000	0.163200	1.000000	0.000000	0.000000	0.000
25%	3.810000	1.500000	3.000000	1.000000	1.000000	0.000
50%	5.940000	2.385600	3.000000	3.000000	2.000000	0.000
75%	8.670000	3.484025	4.000000	4.000000	3.000000	1.000
max	22.920000	9.726500	6.000000	5.000000	4.000000	5.000

Correlation Matrix

Why?

A correlation matrix is a table showing correlation coefficients between variables.

There are three broad reasons for computing a correlation matrix:

To summarize a large amount of data where the goal is to see patterns. In our example above, the observable pattern is that all the variables highly correlate with each other. To input into other analyses. For example, people commonly use correlation matrixes as inputs for exploratory factor analysis, confirmatory factor analysis, structural equation models, and linear regression when excluding missing values pairwise. As a diagnostic when checking other analyses. For example, with linear regression, a high amount of correlations suggests that the linear regression estimates will be unreliable.

```
In [5]: #Correlation of the dataset
data.corr().style.background_gradient(cmap='afmhot')
```

Out[5]:


	store_sales(in millions)	store_cost(in millions)	unit_sales(in millions)	total_children	avg_cars_at home(approx)	nu
store_sales(in millions)	1.000000	0.954685	0.503482	0.083313	0.004498	
store_cost(in millions)	0.954685	1.000000	0.480087	0.079058	0.002865	
unit_sales(in millions)	0.503482	0.480087	1.000000	0.163188	0.023667	
total_children	0.083313	0.079058	0.163188	1.000000	0.098110	
avg_cars_at home(approx)	0.004498	0.002865	0.023667	0.098110	1.000000	
num_children_at_home	0.032437	0.027576	0.066725	0.394709	0.130841	
avg_cars_at home(approx).1	0.004498	0.002865	0.023667	0.098110	1.000000	
SRP	0.833478	0.795880	-0.002358	0.000545	-0.007921	
gross_weight	0.036179	0.034237	0.001255	-0.000186	0.004588	
net_weight	0.032014	0.030257	0.001137	0.000142	0.004155	
recyclable_package	0.034293	0.030213	0.001599	0.002794	0.003725	
low_fat	-0.006134	-0.005976	-0.001129	-0.002824	-0.004312	
units_per_case	-0.010630	-0.009792	0.000084	0.002307	-0.007265	
store_sqft	0.015543	0.017877	0.031464	0.000555	-0.015815	
grocery_sqft	0.010442	0.012884	0.024857	0.018526	-0.017694	
frozen_sqft	0.017886	0.019245	0.030563	-0.026926	-0.007470	
meat_sqft	0.017883	0.019242	0.030557	-0.026923	-0.007466	
coffee_bar	-0.029368	-0.027126	-0.057633	0.002836	-0.002702	
video_store	0.019179	0.019252	0.034996	-0.000591	0.014001	
salad_bar	0.031459	0.033206	0.057878	-0.013764	-0.008982	
prepared_food	0.031459	0.033206	0.057878	-0.013764	-0.008982	
florist	0.030603	0.030929	0.055885	-0.003361	-0.004138	
cost	-0.004621	-0.004162	-0.015015	-0.003900	0.011658	

```

In [6]: # Correlation matrix using seaborn as heatmap
plt.figure(figsize=(19,14))
sns.heatmap(data.corr(),cmap='rainbow',annot=True)
plt.title("The correlation matrix of the dataset",fontsize=40)
plt.show()

```

The correlation matrix of the dataset



EDA PROCESS

```
In [7]: #Dataset columns
data.columns
```

```
Out[7]: Index(['food_category', 'food_department', 'food_family',
        'store_sales(in millions)', 'store_cost(in millions)',
        'unit_sales(in millions)', 'promotion_name', 'sales_country',
        'marital_status', 'gender', 'total_children', 'education',
        'member_card', 'occupation', 'houseowner', 'avg_cars_at home(approx)',
        'avg. yearly_income', 'num_children_at_home',
        'avg_cars_at home(approx).1', 'brand_name', 'SRP', 'gross_weight',
        'net_weight', 'recyclable_package', 'low_fat', 'units_per_case',
        'store_type', 'store_city', 'store_state', 'store_sqft', 'grocery_sqf
t',
        'frozen_sqft', 'meat_sqft', 'coffee_bar', 'video_store', 'salad_bar',
        'prepared_food', 'florist', 'media_type', 'cost'],
        dtype='object')
```

A pie chart is a circular statistical chart, which is divided into sectors to illustrate numerical

proportion.


If you're looking instead for a multilevel hierarchical pie-like chart, go to the Sunburst tutorial.

Pie chart with plotly express

To visualize the food_category in dataset

```
In [8]: #Let's visualize the food_category in the dataset using the plotly
fig=px.pie(data,names='food_category',title='To Visualize the food_category in
    'Vegetables', 'Frozen Desserts', 'Candy', 'Snack Foods', 'Dairy',
    'Starchy Foods', 'Cleaning Supplies', 'Decongestants', 'Meat',
    'Hot Beverages', 'Jams and Jellies', 'Carbonated Beverages',
    'Seafood', 'Specialty', 'Kitchen Products', 'Electrical',
    'Beer and Wine', 'Candles', 'Fruit', 'Pure Juice Beverages',
    'Canned Soup', 'Paper Products', 'Canned Tuna', 'Eggs', 'Hardware',
    'Canned Sardines', 'Canned Clams', 'Pain Relievers', 'Side Dishes',
    'Bathroom Products', 'Magazines', 'Frozen Entrees', 'Pizza',
    'Cold Remedies', 'Canned Anchovies', 'Drinks', 'Hygiene',
    'Plastic Products', 'Canned Oysters', 'Packaged Vegetables',
    'Miscellaneous'])
fig.update_traces(textposition='inside')
fig.update_layout(uniformtext_minsize=12, uniformtext_mode='hide')
fig.show()
```

To Visualize the food_category in dataset




Observation:

From the above pie chart most used food_category such as 1.Vegetables 2.Snack Foods 3.Dairy these products we used in our daily life

To visualize the food_department in the dataset

```
In [9]: #Let's create a pie chart using the plotly to visualize the food_department in
fig = go.Figure(data=[go.Pie(labels=['Frozen Foods', 'Baked Goods', 'Canned Fo
    'Produce', 'Snacks', 'Snack Foods', 'Dairy', 'Starchy Foods',
    'Household', 'Health and Hygiene', 'Meat', 'Beverages', 'Seafood',
    'Deli', 'Alcoholic Beverages', 'Canned Products', 'Eggs',
    'Periodicals', 'Breakfast Foods', 'Checkout', 'Carousel'],values=data['
#update the piechart with colors and text info and border with line
fig.update_traces(hoverinfo='label+percent', textinfo='value', textfont_size=2
    marker=dict(colors=['Crimson','Chocolate','mediumturquoise',
fig.show()
```


Observation:

From the pie chart the most used food_department in the dataset such as 1.Frozen Foods, 2.Baked Goods, 3.Canned Foods.

To Visualize the sales_country in the dataset

```
In [10]: #Let's create a pie chart to visualize the sales_country in the dataset with t
fig=px.pie(data,names='sales_country',title='To Visualize the sales_country in
#update the pie chart with the colors and border line and finally visualized
fig.update_traces(textposition='inside')
fig.update_layout(uniformtext_minsize=12, uniformtext_mode='hide')
fig.show()
```

To Visualize the sales_country in dataset



Observation:

From the above pie chart most sales done in the USA, After that Mexico and finally least sales done in the Canada

Count Plot

A countplot is kind of like a histogram or a bar graph for some categorical area.

It simply shows the number of occurrences of an item based on a certain type of category.

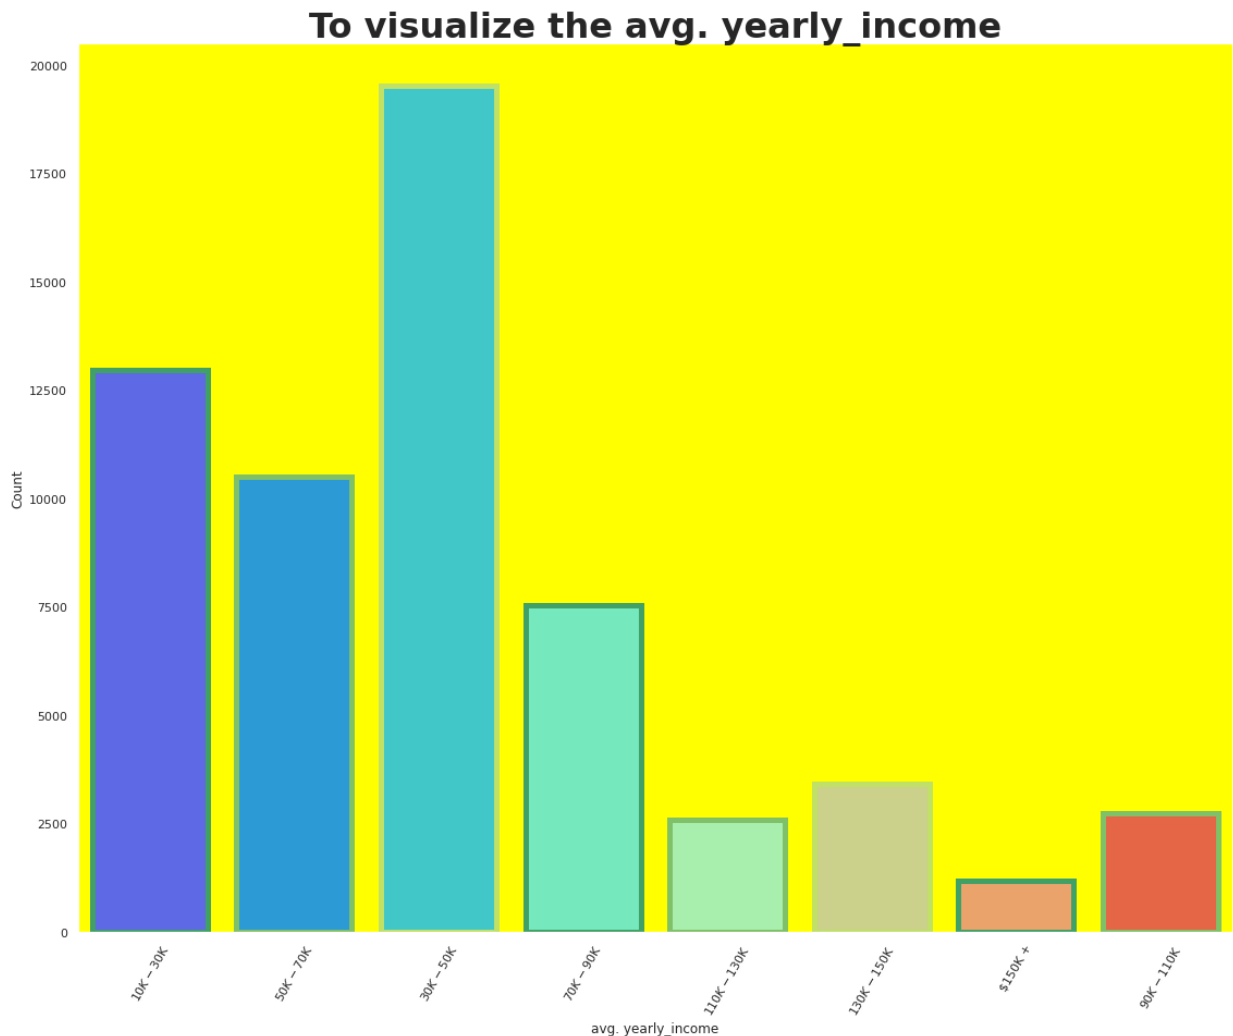
```
In [11]: # count plot of whole dataset based on promotion_name
ax=plt.axes()
#Set the background color
ax.set(facecolor='black')
#set the figure size and style
sns.set(rc={'figure.figsize':(19,15)},style='dark')
#create the title of the plot
ax.set_title("To visualize the promotion_names",fontsize=32,fontweight='bold')
#create the countplot using the seaborn with the params
```

```
sns.countplot(data[promotion_name'],palette='rainbow',linewidth=5,edgecolor=s
#on the x-axis the promotion_names
plt.xlabel('promotion_name')
#on the y_axis the count of the promotion
plt.ylabel('Count')
#creat the ticks on x axis beacause to visualize the botom promotion_names
plt.xticks(rotation=60)
#finaly visualize it
plt.show()
```

```

ax.set_title("To visualize the avg. yearly_income",fontsize=32,fontweight='bold')
#create the countplot using the seaborn with the params
sns.countplot(data['avg. yearly_income'],palette='rainbow',linewidth=5,edgecolor='black')
#on the x-axis the promotion_names
plt.xlabel('avg. yearly_income')
#on the y-axis the count of the promotion
plt.ylabel('Count')
#creat the ticks on x axis beacause to visualize the botom avg. yearly_income
plt.xticks(rotation=60)
#finally visualize it
plt.show()

```



1) The avg. yearly_income 30k - 50 k

Histograms


In statistics, a histogram is representation of the distribution of numerical data, where the data are binned and the count for each bin is represented. More generally, in Plotly a histogram is an aggregated bar chart, with several possible aggregation functions (e.g. sum, average, count...) which can be used to visualize data on categorical and date axes as well as linear axes.

To visualize the Education with gender in the

dataset using the histogram


```
In [13]: #Let's create the histogram to visualize the To visualize the education with g
#Create the histogram to visuzlized
fig=px.histogram(data,x='education',color='gender',title='To visualize the edu
fig.update_layout(
    title_text='Sampled Results', # title of plo
    bargap=0.2, # gap between bars of adjacent location coordinates
    bargroupgap=0.1, # gap between bars of the same location coordinates
    plot_bgcolor='pink'
)
fig.show()
```

Sampled Results




```
In [14]: #Let's create the histogram to visualize the To visualize the occupation with
#Create the histogram to visuzlized
fig=px.histogram(data,x='occupation',color='gender',title='To visualize the ed
fig.update_layout(bargap=0.2,bargroupgap=0.1,
    plot_bgcolor='green'
)
fig.show()
```

To visualize the education with gender



```
In [15]: #Let's create the histogram to visualize the To visualize the occupation with
#Create the histogram to visuzlized
fig=px.histogram(data,color='occupation',x='brand_name',title='To visualize th
fig.update_layout(bargap=0.2,bargroupgap=0.1,
                  plot_bgcolor='green'
)
fig.show()
```

To visualize the brand_name with occupation



Observation:

For Male:


1. Most Male person have the partial high scholl and high school degree 2. Les number have males have Graduate degree

For Female:

1. Most FeMale person have the partial high scholl and high school degree 2. Les number have Females have Graduate degree


```
In [16]: #to visaulize the occupation with education using the bar plotly express
fig = px.histogram(data, color='avg. yearly_income', x='member_card',title='To
fig.update_layout(bargap=0.2,bargroupgap=0.1,plot_bgcolor='orange')
fig.show()
```

To visualize the avg. yearly_income with member_card




```
In [17]: #to visualize the occupation with education using the bar plotly express
fig = px.histogram(data, color='occupation', x='education', title='To visualize
fig.update_layout(bargap=0.2, bargroupgap=0.1, plot_bgcolor='Navy')
fig.show()
```


To visualize the education with occupation



```
In [18]: #to visualize the occupation with avg. yearly_income using the bar plotly expr
fig = px.histogram(data, color='occupation', x='avg. yearly_income',title='To
fig.update_layout(bargap=0.2,bargroupgap=0.1,plot_bgcolor='Green')
fig.show()
```

To visualize the avg. yearly_income with occupation



Information:

From the Professional employee average earn 50K-


From the skilled Manual Employee and Manual Employee 30 K-

From the Management Employee earn 70 K-

From Clerical Employee Earn 30K-


```
In [19]: #to visaulize the store_sales(in millions) with store_city using the bar plotl
fig = px.histogram(data, x='store_sales(in millions)', color='store_city', titl
fig.update_layout(bargap=0.2,bargroupgap=0.1,plot_bgcolor='orange')
fig.show()
```

To visualize the avg. yearly_income with occupation



```
In [20]: #to visualize the store_sales(in millions) with store_state using the bar plot
fig = px.histogram(data, color='store_type', x='store_state',title='To visuali
fig.update_layout(bargap=0.2,bargroupgap=0.1,plot_bgcolor='maroon')
fig.show()
```

To visualize the store_type with store_state




DISTPLOT

The distplot represents the univariate distribution of data i.e. data distribution of a variable against the density distribution

```
In [21]: #import the norm from scipy
from scipy.stats import norm
#Visualize the distplot
sns.distplot(data['cost'], fit=norm, kde=False)
```


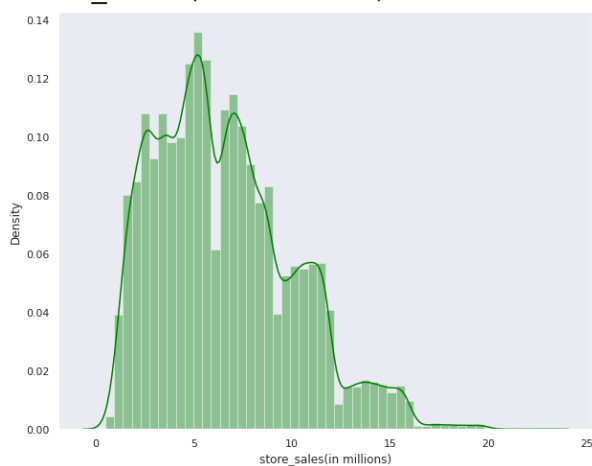
```
Out[21]: <AxesSubplot: xlabel='cost'>
```



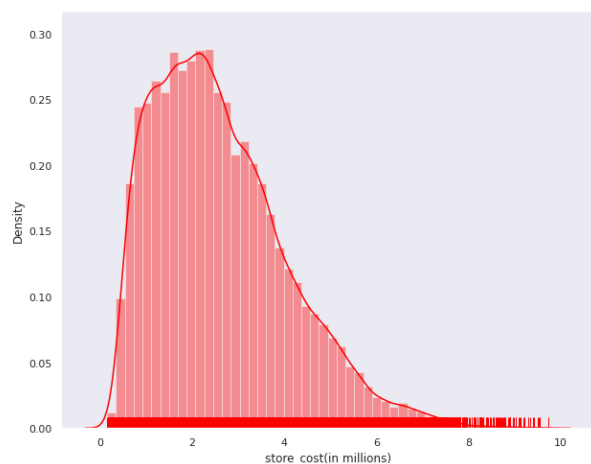
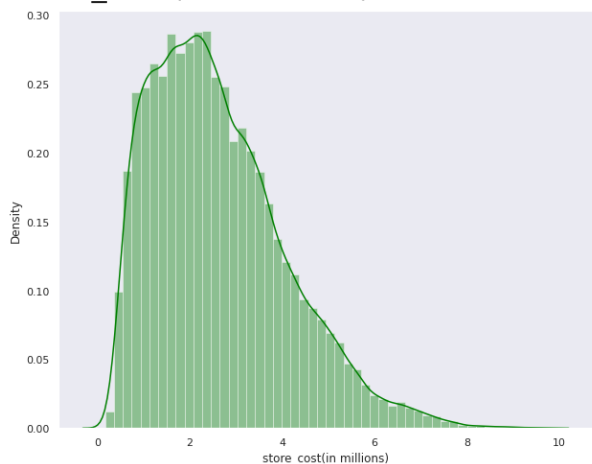
```
In [22]: #Perform distplot for all the columns in dataset
col=data[['store_sales(in millions)', 'store_cost(in millions)', 'unit_sales(in
net_weight', 'recyclable_package', 'low_fat', 'units_per_case', 'store_sqf
for column in col.columns:
    print(column)
    # code below
    fig,ax = plt.subplots(nrows=1,ncols=2,figsize=(22,8))
    sns.distplot(data[column],ax=ax[0],color='green',hist=True)
    sns.distplot(data[column],ax=ax[1],rug=True,hist=True,color='red')

    plt.show()
```



store_sales(in millions)





store_cost(in millions)




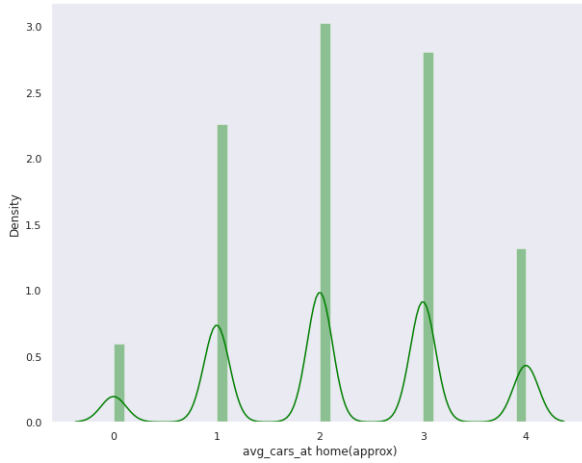
unit_sales(in millions)





total_children





avg_cars_at home(approx)





num_children_at_home





avg_cars_at home(approx).1





SRP





gross_weight





net_weight





recyclable_package





low_fat





units_per_case





store_sqft





grocery_sqft





frozen_sqft





meat_sqft





coffee_bar




video_store



salad_bar




prepared_food




```
In [23]: # violin plot for store_sales(in millions) and store_cost(in millions) columns
fig=px.violin(data,x='store_sales(in millions)', y='store_cost(in millions)',t
fig.show()
```

The relation between store_sales(in millions) and store_cost(in millions)



```
In [24]: # violin plot for store_sqft and grocery_sqft columns
fig=px.scatter(data,x='store_sqft', y='grocery_sqft',title="The relation betwe
fig.update_layout(bargap=0.2,bargroupgap=0.1,plot_bgcolor='red')
fig.show()
```

The relation between store_sqft and grocery_sqft



Point Plot


A point plot uses scatter plot glyphs to visualize features like point estimates and confidence intervals.

A point plot uses scatter plot points to represent the central tendency of numeric data.

These plots make use of error bars to indicate any uncertainty around the numeric


```
In [25]: # point plot for milesFromMetropolis from salary columns
plt.figure(figsize=(20,6))
sns.pointplot(x='frozen_sqft', y='meat_sqft', data=data, palette='rainbow')
plt.title("The relation meat_sqft and frozen_sqft", fontsize=32)
plt.show()
```

The relation meat_sqft and frozen_sqft




```
In [26]: #To visualize the pointplot grocery_sqft and store_sqft
plt.figure(figsize=(20,6))
sns.pointplot(x='store_sqft', y='grocery_sqft', data=data, palette='rainbow')
plt.title("The relation store_sqft and grocery_sqft",fontsize=32)
plt.show()
```

The relation store_sqft and grocery_sqft




```
In [27]: #To visualize the pointplot gross_weight and net_weight
plt.figure(figsize=(20,6))
sns.pointplot(x='gross_weight', y='net_weight', data=data, palette='rainbow')
plt.title("The relation gross_weight and net_weight",fontsize=32)
plt.show()
```

The relation gross_weight and net_weight



```
In [28]: #To visualize the pointplot gross_weight and units_per_case
plt.figure(figsize=(20,6))
sns.pointplot(x='gross_weight', y='units_per_case', data=data, palette='rainbow')
plt.title("The relation recyclable_package and low_fat",fontsize=32)
plt.show()
```




Strip Plot

A strip plot is a graphical data analysis technique for summarizing a univariate data set. The strip plot consists of:

1. Horizontal axis = the value of the response variable;
2. Vertical axis = all values are set to 1.


That is, a strip plot is simply a plot of the sorted response values along one axis. The strip plot is an alternative to a histogram or a density plot. It is typically used for small data sets (histograms and density plots are typically preferred for larger data sets).

```
In [29]: # point stripplot for milesFromMetropolis from salary columns
plt.figure(figsize=(20,6))
sns.stripplot(x='frozen_sqft', y='meat_sqft', data=data, palette='rainbow')
plt.title("The relation meat_sqft and frozen_sqft", fontsize=32)
plt.show()
```




```
In [30]: #To stripplot the pointplot grocery_sqft and store_sqft
plt.figure(figsize=(20,6))
sns.stripplot(x='store_sqft', y='grocery_sqft', data=data, palette='rainbow')
plt.title("The relation store_sqft and grocery_sqft", fontsize=32)
plt.show()
```

The relation store_sqft and grocery_sqft



```
In [31]: #To visualize the pointplot gross_weight and net_weight
plt.figure(figsize=(20,6))
sns.stripplot(x='gross_weight', y='net_weight', data=data, palette='rainbow')
plt.title("The relation gross_weight and net_weight", fontsize=32)
plt.show()
```

The relation gross_weight and net_weight



```
In [32]: #Check the columns
data.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 60428 entries, 0 to 60427
Data columns (total 40 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   food_category                             60428 non-null  object
1   food_department                           60428 non-null  object
2   food_family                               60428 non-null  object
3   store_sales(in millions)                  60428 non-null  float64
4   store_cost(in millions)                   60428 non-null  float64
5   unit_sales(in millions)                   60428 non-null  float64
6   promotion_name                            60428 non-null  object
7   sales_country                             60428 non-null  object
8   marital_status                           60428 non-null  object
9   gender                                    60428 non-null  object
10  total_children                            60428 non-null  float64
11  education                                 60428 non-null  object
12  member_card                              60428 non-null  object
13  occupation                               60428 non-null  object
14  houseowner                               60428 non-null  object
15  avg_cars_at home(approx)                  60428 non-null  float64
16  avg. yearly_income                        60428 non-null  object
17  num_children_at_home                      60428 non-null  float64
18  avg_cars_at home(approx).1                60428 non-null  float64
19  brand_name                                60428 non-null  object
20  SRP                                       60428 non-null  float64
21  gross_weight                             60428 non-null  float64
22  net_weight                               60428 non-null  float64
23  recyclable_package                       60428 non-null  float64
24  low_fat                                  60428 non-null  float64
25  units_per_case                           60428 non-null  float64
26  store_type                               60428 non-null  object
27  store_city                               60428 non-null  object
28  store_state                              60428 non-null  object
29  store_sqft                               60428 non-null  float64
30  grocery_sqft                             60428 non-null  float64
31  frozen_sqft                             60428 non-null  float64
32  meat_sqft                               60428 non-null  float64
33  coffee_bar                               60428 non-null  float64
34  video_store                              60428 non-null  float64
35  salad_bar                                60428 non-null  float64
36  prepared_food                            60428 non-null  float64
37  florist                                  60428 non-null  float64
38  media_type                               60428 non-null  object
39  cost                                     60428 non-null  float64
dtypes: float64(23), object(17)
memory usage: 18.4+ MB
```

Modeling

```
In [33]: #Convert the categorical columns to numerical using the LabelEncoder
from sklearn.preprocessing import LabelEncoder
label=LabelEncoder()
data['food_category']=label.fit_transform(data['food_category'])
data['food_department']=label.fit_transform(data['food_department'])
data['food_family']=label.fit_transform(data['food_family'])
data['promotion_name']=label.fit_transform(data['promotion_name'])
data['sales_country']=label.fit_transform(data['sales_country'])
```

```
data['marital_status']=label.fit_transform(data['marital_status'])
data['gender']=label.fit_transform(data['gender'])
data['education']=label.fit_transform(data['education'])
data['member_card']=label.fit_transform(data['member_card'])
data['occupation']=label.fit_transform(data['occupation'])
data['avg. yearly_income']=label.fit_transform(data['avg. yearly_income'])
data['brand_name']=label.fit_transform(data['brand_name'])
data['houseowner']=label.fit_transform(data['houseowner'])
data['store_type']=label.fit_transform(data['store_type'])
data['store_city']=label.fit_transform(data['store_city'])
data['store_state']=label.fit_transform(data['store_state'])
data['media_type']=label.fit_transform(data['media_type'])
```

```
In [34]: #Divided the dataset into x and y variables
X=data.iloc[:, :-1]
y=data['cost']
```

```
In [35]: # Helper function for scaling all the numerical data using MinMaxScaler
# import asarray
# import MinMaxScaler
#def scale_data(df,col):
#    scaler = MinMaxScaler()
#    df[col] = scaler.fit_transform(df[col])

#    return df
```

```
In [36]: #col=['store_sales(in millions)', 'store_cost(in millions)', 'unit_sales(in mill
#         'net_weight', 'recyclable_package', 'low_fat', 'units_per_case', 'store
#         'frozen_sqft', 'meat_sqft', 'coffee_bar', 'video_store', 'salad_bar',
#         'prepared_food']
#X = scale_data(data,col)
```

```
In [37]: #Divided the data into train and test data
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,random_state
```

```
In [38]: #Let's print the train and test shape
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(45321, 39)
(15107, 39)
(45321,)
(15107,)
```

LinearRegression model

```
In [39]: #Install the LinearRegression model to predict the cost
linear=LinearRegression()
#fit the model to the train data
linear.fit(X_train,y_train)
```

```
Out[39]: LinearRegression()
```

```
In [40]: #Prediction of the LinearRegression mode
```

```
linear_pred=linear.predict(X_test)
linear_pred
```

```
Out[40]: array([ 97.60471773, 104.94182506, 93.09602181, ..., 98.16966396,
        96.30744502, 94.56495128])
```

```
In [41]: #Check the test score and train score to the LinearRegression algorithm
print(f'The Test_accuracy: {linear.score(X_test,y_test)*100:.2f}')
#Train score for the data
print(f'The Train_accuracy: {linear.score(X_train,y_train)*100:.2f}')
```

```
The Test_accuracy: 3.55
The Train_accuracy: 3.35
```

Mean_squared_error and r2_score to the linearRegression model

```
In [42]: #Install the mean_squared_error and r2_score from the sklearn libraries
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
mse=mean_squared_error(y_test, linear_pred)
rmse=np.sqrt(mse)
print("Root_mean_squred_error LinearRegression {:.4f}".format(rmse))
print("R2_score LinearRegression {:.4f}".format(r2_score(y_test, linear_pred)))
print("mean_absolute_error LinearRegression {:.4f}".format(mean_absolute_error(
```

```
Root_mean_squred_error LinearRegression 29.5271
R2_score LinearRegression 0.035500
mean_absolute_error LinearRegression 25.422579
```

DecisionTreeRegressor

```
In [43]: #Import the DecisionTreeRegressor from the sklearn
from sklearn.tree import DecisionTreeRegressor
#Install the DecisionTreeRegressor model
tree=DecisionTreeRegressor()
#And finally we fit the train and test the data
tree.fit(X_train,y_train)
```

```
Out[43]: DecisionTreeRegressor()
```

```
In [44]: #Check the test score and train score to the DecisionTreeRegressor algorithm
print(f'The Test_accuracy: {tree.score(X_test,y_test)*100:.2f}')
#Train score for the data
print(f'The Train_accuracy: {tree.score(X_train,y_train)*100:.2f}')
```

```
The Test_accuracy: 99.75
The Train_accuracy: 100.00
```

```
In [45]: #Predictio ot the DecisionTreeRegressor
tree_pred=tree.predict(X_test)
tree_pred
```

```
Out[45]: array([111.7 , 145.6 , 92.57, ..., 57.52, 62.74, 99.38])
```

Mean_squared_error and r2_score to the linearRegression model

```
In [46]: mse=mean_squared_error(y_test,tree_pred)
rmse=np.sqrt(mse)
print("Root_mean_squred_error DecisionTreeRegressor {:.4f}".format(rmse))
print("R2_score DecisionTreeRegressor {:.4f}".format(r2_score(y_test,tree_pred))
print("mean_absolute_error DecisionTreeRegressor {:.4f}".format(mean_absolute_e
```

```
Root_mean_squred_error DecisionTreeRegressor 1.4910
R2_score DecisionTreeRegressor 0.997541
mean_absolute_error DecisionTreeRegressor 0.050976
```

RandomForestRegressor

```
In [47]: #Install the RandomForestRegressor from the the sklearn
from sklearn.ensemble import RandomForestRegressor
#Install the RandomForestRegressor model
random=RandomForestRegressor()
#Fit the train data to the model
random.fit(X_train,y_train)
```

```
Out[47]: RandomForestRegressor()
```

```
In [48]: #Predictionof the RandomForestRegressor algorithm
random_pred=random.predict(X_test)
random_pred
```

```
Out[48]: array([111.7 , 145.6 , 92.57, ..., 57.52, 62.74, 99.38])
```

```
In [49]: #Check the test score and train score to the RandomForestRegressor algorithm
print(f'The Test_accuracy: {random.score(X_test,y_test)*100:.2f}')
#Train score for the data
print(f'The Train_accuracy: {random.score(X_train,y_train)*100:.2f}')
```

```
The Test_accuracy: 99.88
The Train_accuracy: 99.98
```

Mean_squared_error and r2_score to the linearRegression model

```
In [50]: #RandomForestRegressor algorithms mean_squared_error and r2_score
mse=mean_squared_error(y_test,random_pred)
rmse=np.sqrt(mse)
print("Root_mean_squred_error RandomForestRegressor {:.4f}".format(rmse))
print("R2_score RandomForestRegressor {:.4f}".format(r2_score(y_test,random_pre
print("mean_absolute_error RandomForestRegressor {:.4f}".format(mean_absolute_e
```

```
Root_mean_squred_error RandomForestRegressor 1.0556
R2_score RandomForestRegressor 0.998767
mean_absolute_error RandomForestRegressor 0.080142
```

XGBRegressor

```
In [51]: # Import XGBRegressor
from xgboost import XGBRegressor
# Instantiate the model
xgb=XGBRegressor()
# Fit the model to the data
# Fit the model to the data
xgb.fit(X_train,y_train)
```

```
Out[51]: XGBRegressor(base_score=0.5, booster=None, colsample_bylevel=1,
                    colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
                    importance_type='gain', interaction_constraints=None,
                    learning_rate=0.300000012, max_delta_step=0, max_depth=6,
                    min_child_weight=1, missing=nan, monotone_constraints=None,
                    n_estimators=100, n_jobs=0, num_parallel_tree=1, random_state=0,
                    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
                    tree_method=None, validate_parameters=False, verbosity=None)
```

```
In [52]: #Prediction to the XGBRegressor algorithms
xgb_pred=xgb.predict(X_test)
xgb_pred
```

```
Out[52]: array([111.88061 , 145.49605 ,  92.479965, ...,  57.674427,  62.387634,
                99.29608 ], dtype=float32)
```

```
In [53]: #Check the test score and train score to the XGBRegressor algorithm
print(f'The Test_accuracy: {xgb.score(X_test,y_test)*100:.2f}')
#Train score for the data
print(f'The Train_accuracy: {xgb.score(X_train,y_train)*100:.2f}')
```

The Test_accuracy: 99.83
The Train_accuracy: 99.88

Mean_squared_error and r2_score to the linearRegression model

```
In [54]: #XGBRegressor algorithms mean_squared_error and r2_score
mse=mean_squared_error(y_test,xgb_pred)
rmse=np.sqrt(mse)
print("Root_mean_squared_error XGBRegressor {:.4f}".format(rmse))
print("R2_score XGBRegressor {:.4f}".format(r2_score(y_test,xgb_pred)))
print("mean_absolute_error XGBRegressor {:.4f}".format(mean_absolute_error(y_te
```

Root_mean_squared_error XGBRegressor 1.2226
R2_score XGBRegressor 0.998346
mean_absolute_error XGBRegressor 0.447202

LGBMRegressor

```
In [55]: # Import LGBMRegressor
from lightgbm import LGBMRegressor
# Instantiate the model
lgb=LGBMRegressor()
```

```
# Fit the model to the data
lgb.fit(X_train,y_train)
```

Out[55]: LGBMRegressor()

```
In [56]: #Prediction of the LGBMRegressor algorithms
lgb_pred=lgb.predict(X_test)
lgb_pred
```

Out[56]: array([107.43654254, 135.94695349, 92.43485135, ..., 61.75415069,
64.79114351, 99.25508362])

```
In [57]: #Check the test score and train score to the LGBMRegressor algorithm
print(f'The Test_accuracy: {xgb.score(X_test,y_test)*100:.2f}')
```

```
#Train score for the data
print(f'The Train_accuracy: {xgb.score(X_train,y_train)*100:.2f}')
```

The Test_accuracy: 99.83
The Train_accuracy: 99.88

Mean_squared_error and r2_score to the linearRegression model

```
In [58]: #LGBMRegressor algorithms mean_squared_error and r2_score
mse=mean_squared_error(y_test,lgb_pred)
rmse=np.sqrt(mse)
print("Root_mean_squred_error LGBMRegressor {:.4f}".format(rmse))
print("R2_score LGBMRegressor {:.4f}".format(r2_score(y_test,lgb_pred)))
print("mean_absolute_error LGBMRegressor {:.4f}".format(mean_absolute_error(y_t
```

Root_mean_squred_error LGBMRegressor 4.5399
R2_score LGBMRegressor 0.977199
mean_absolute_error LGBMRegressor 3.237443

From the Above data the DecisionTreeRegressor, RandomForestRegressor,XGBRegressor and LGBRegressor give the better accuracy_score

About the dataset

The data information is taken from the kaggle website. This data data preprocessing involves several steps, firstly we do the basic of the process, And then do some EDA process we visualize the pie chart, histogram, distplot, stripplot, pointplot and then we convert the categorical data to numerical data using the label encoder. After that we divided the modeling process after that spilt the train and test and we ready to the modelig. we use several algorithms to predict the output such as LinearRegressor, DecisionTreeRegressor, RandomForestRegressor,XGBRegressor and LGBRegressor.

In []: